



香港中文大學(深圳)
The Chinese University of Hong Kong, Shenzhen

Weather sensitivity of bike-share usage a hierarchical Bayesian plan for 25 cities

STA5001 Bayesian & Computational Statistics

XUE Yantong

student1@example.com

WANG Ziyu

student2@example.com

ZHAO Jiashen

student3@example.com

STA5001 · 3 November 2026

Weather drives bike-share usage

The reference study. Bean, Pojani and Corcoran (2021) compare **forty cities** across climate zones, fitting a separate generalized additive model to each.

Two robust patterns: hires rise with UTCI up to about **27–28°C** and then fall; rainfall is associated with lower use.

What a per-city fit cannot answer

- How much cities differ — no quantity describes the spread.
- Small systems are estimated noisily.

Question. Do a city's thermal and rainfall background explain part of its weather sensitivity?

Data: one public release, 29 reproducible cities

Source

Release	Mendeley Data, CC BY 4.0
DOI	10.17632/2nxpvtz935.1
Content	hourly hires, UTCI, rainfall
Scope	29 dock-based systems

Why 29 and not 40. The paper's own code designates the dock-based systems; the other cities have metadata but no hourly stock data.

Each city contributes exactly **8,760** hourly rows.

Archive and paper are both stored with SHA-256 digests, re-checked on every run.

From hourly records to 9,120 analysis days

Processing

- Preserve row order, map onto consecutive UTC hours
- Convert to the city's local time zone
- Aggregate by local date


Quality rules

- 23–25 hourly rows per local day
- at least 20 valid UTCI hours
- complete precipitation

254,040 hourly rows → 10,613 candidate city-days → **10,578** kept → **9,120** in the primary sample of 25 cities


Three repeated time keys exist; Dublin and Melbourne coincide with daylight-saving transitions, Brisbane does not observe daylight saving; the paper's row-order rule is kept.

Cities differ in scale by four orders of magnitude



Median daily hires run from **11.5** in Créteil to **76,103** in Paris.
The four cities below the 99% line are excluded from the primary sample.


Thermal response peaks near 27.5 °C



Peaks at **+0.26** near **27.5 °C**, then falls to **-0.14** in the hottest bin.

Rainfall declines monotonically: the heaviest bin sits about **41%** below the dry bin.

Season and week disagree across cities



The seasonal amplitude of **0.84** log units exceeds the whole thermal effect.

23 of 25 cities have lower weekend use; Melbourne is the only clearly positive city.

Plan: one hierarchy instead of 25 separate fits

Response. $y_{it} = \log(1 + U_{it})$, Student- t observation distribution with estimated degrees of freedom.

Location. city intercept, thermal anomaly and its square, standardised rainfall, weekend, annual sine and cosine.

Hierarchy. city coefficients share a normal population distribution, written in non-centred form.

What this buys

- Between-city spread is an estimated quantity
- Population means vary with city climate background
- Low-volume systems borrow strength

NUTS via PyMC and nutpie, four chains of 2,000 draws.

Take-away

1. A reproducible panel of **25 cities and 9,120 city-days** is frozen and documented.
2. The thermal response is **non-monotone**, peaking near **27.5°C**, and rainfall is associated with monotonically lower use.
3. Cities differ in scale, in weekend behaviour and in their weather response — which is why the final model is a **hierarchy**.

Questions

Backup: filtering and sample construction

Stage	Result
Hourly stock rows for 29 cities	254,040
Candidate local days	10,613
Days failing a quality rule	35
Retained city-days	10,578
Cities recording hires on $\geq 99\%$ of days	25
Primary sample	9,120 city-days

Exclusions overlap: 28 days outside the nominal study period, 34 outside the 23–25 hour count, 33 short of 20 UTCI hours, none on precipitation. The four excluded cities are Göteborg, Kazan, Lillestrøm and Vilnius.

Backup: model and priors

Priors on the $\log(1 + U)$ scale

Population intercept	$\mathcal{N}(7, 2)$
City intercept scale	$\text{HN}(2)$
Thermal, per 10 °C	$\mathcal{N}(0, 0.5)$
Quadratic	$\mathcal{N}(0, 0.3)$
Rainfall	$\mathcal{N}(0, 0.3)$
Weekend	$\mathcal{N}(0, 0.4)$
Seasonal	$\mathcal{N}(0, 0.8)$


Degrees of freedom $\nu = 2 + \text{Exponential}(0.1)$; moderator coefficients use $\mathcal{N}(0, 0.2)$ to $\mathcal{N}(0, 0.3)$, and city scales use HalfNormal priors matched to the population scale.

A prior predictive check will verify the implied range of daily hires and weather multipliers before the observed outcomes are examined.


Convergence thresholds: $\hat{R} \leq 1.01$, bulk and tail ESS at 400, no divergences.

Backup: per-city curves for six representative cities

Thermal response, six representative cities



Rainfall response, six representative cities



Lund (1.2 °C), Amiens (6.5), Toyama (10.9), Valence (16.3), Brisbane (19.9) and Kaohsiung (26.6) are equally spaced over the range of city means; the grids for all 25 cities are in the supplementary figures.

Backup: source and reproducibility

- Data release: Mendeley Data, DOI 10.17632/2nxpvtz935.1, CC BY 4.0, Version 1.
- Reference paper: Bean, Pojani and Corcoran (2021), *Journal of Transport Geography* 95:103155.
- Preparation: `python code/prepare_bikeshare_weather.py`
- Daily panel: `data/clean/bikeshare_weather_daily.csv`, 10,578 rows $\times$ 28 columns.
- Pilot model: `code/run_pilot_model.sh`, PyMC 6.3.2 with nutpie. No divergences, but several city-level parameters mixed poorly and some transitions hit the maximum tree depth, so it is a feasibility test.
- Verification: `python code/verify_project_setup.py`